

SciTec MLOps Engineering Coding Challenge Instructions

Contents

1. Problem Statement
2. Context
3. Input Data
4. Requirements
5. Allowed Resources
6. What to Deliver
7. Assessment Criteria
8. Test Environment

1. Problem Statement

You are tasked with designing and demonstrating a machine learning solution for classifying missile flight phases. Your implementation will involve building an LSTM model for classification, deploying it in a containerized environment, and following MLOps best practices for operationalization. The solution should emphasize scalability, robustness, and adherence to modern MLOps standards.

Your task can be summarized as follows:

- Develop a robust LSTM model to classify missile track segments into reentry and non-reentry phases.
- Deploy your solution using Docker and Kubernetes (Minikube).
- Demonstrate scalability, logging, and monitoring capabilities.

2. Context

Missile Tracking and Classification

- Simulated sensors track ballistic missiles and record telemetry data.
- The dataset includes position (latitude, longitude, altitude), radiometric intensity, and timestamps.
- Your goal is to classify whether a missile is in the reentry phase using labeled training data.

Deployment Requirements

- The model should be deployable in a containerized environment using Kubernetes (Minikube) for orchestration.
- The system must log inference times and predictions for monitoring.

3. Input Data

Files

The following files are provided:

- `train.csv`: Labeled data for training.
 - Columns: `timestamp`, `track_id`, `sensor_id`, `latitude`, `longitude`, `altitude`, `radiometric_intensity`, `reentry_phase`.
- `test.csv`: Unlabeled data for testing.
 - Same columns as `train.csv`, excluding `reentry_phase`.

Notes

- Sensor data includes noise to simulate real-world conditions.
- Preprocessing and feature engineering are encouraged to improve model performance.

4. Requirements

Languages and Tools

- Use Python for implementation.
- Use TensorFlow or PyTorch for model development.

Machine Learning

- Build an LSTM model for classification.

- Preprocess data:
 - Normalize features.
 - Handle noisy data.
- Use regularization techniques (e.g., dropout, batch normalization) and callbacks (e.g., early stopping).
- Evaluate the model against a simple baseline classifier (e.g., altitude threshold).

Deployment and Operations

- Containerize the solution using Docker.
- Deploy the containerized model to Minikube using Kubernetes manifests.
- Include logging to capture inference times and predictions.

Scalability and Testing

- Ensure the solution can scale to handle larger datasets or higher inference loads.
- Provide unit tests and coverage reports to validate functionality.

5. Allowed Resources

- Use any free and open-source software.
- Consult external resources with proper citations.

6. What to Deliver

Code

- Modular Python scripts, structured for production.
- Include unit tests and coverage reporting.

Documentation

- Instructions for running the solution locally and in Minikube.
- A detailed README.md describing:
 - Model architecture and rationale.
 - Data preprocessing and feature engineering steps.
 - Scalability considerations and trade-offs.
 - Comparison to baseline performance.

Deployment

- Dockerfile and Kubernetes manifests for deployment.

- Logs from deployed inference runs.

Performance Validation

- Metrics for model evaluation and inference times.
- Description of scalability testing methodology.

7. Assessment Criteria

Scores for each category will be weighted according to the percentages in **Table 1**. Final scores are a weighted sum of technical skills, decision-making, and creativity.

8. Test Environment

Your solution will be tested on a Linux system with:

- **CPU:** 24 CPUs @ 3.5 GHz
- **RAM:** 64 GB
- **Storage:** 1 TB
- **GPU:** NVIDIA RTX A2000
- **OS:** RHEL 9.5

Table 1. Coding Challenge Evaluation Criteria

Category	Weight	Evaluation Criteria
Model Implementation	30%	Completeness, correctness, and adherence to ML best practices, including preprocessing and evaluation.
Deployment and Ops	30%	Correctness and efficiency of containerization, deployment, and logging.
Code Quality	20%	Modularity, documentation, and adherence to Python best practices.
Performance Evaluation	10%	Use of appropriate metrics and comparison to a baseline.
Scalability and Testing	10%	Demonstration of scalability and robust testing practices.